

Manual for Stacklook

Using Stacklook in KernelShark

How to build

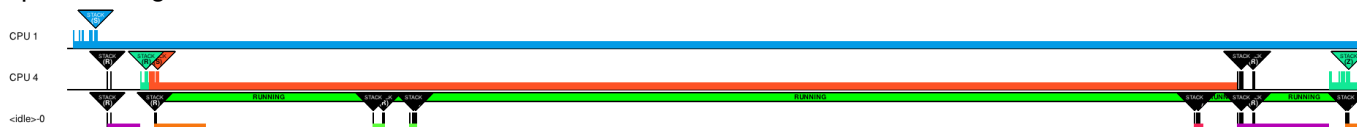
Please follow the [guide in README](#).

Load it up

Having built a KernelShark shared library, remember where it is. Then, when starting KernelShark from the terminal, add `-p [path/to/plugin/binary]` to the `kernelshark` command. This approach will follow symlinks.

Alternatively, you can start KernelShark yourself and then use the GUI the add the plugin that way via the `Tools > Add plugins` button.

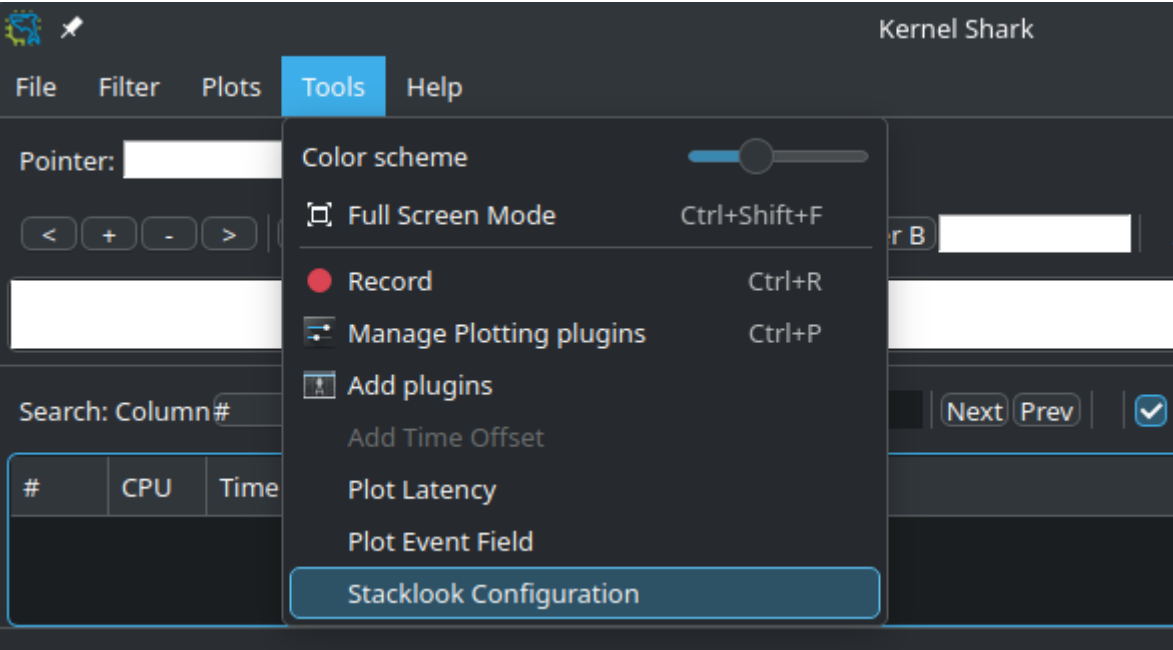
If you have VERY few entries, you might see something like this upon loading:



Configuration

Plugin configuration can be done at any time, even before any trace file is loaded. To use it, simply open the dialog window through "Tools > Stacklook Configuration" in the main window. There can be only one window open at a given time.

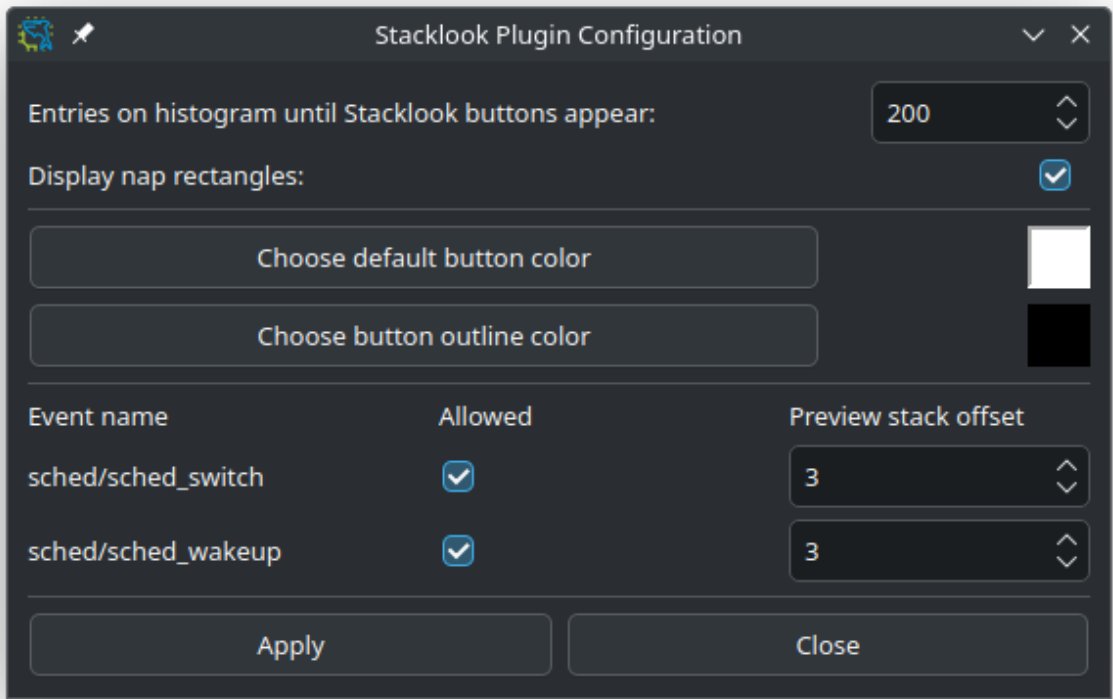
Configuration path as a picture:



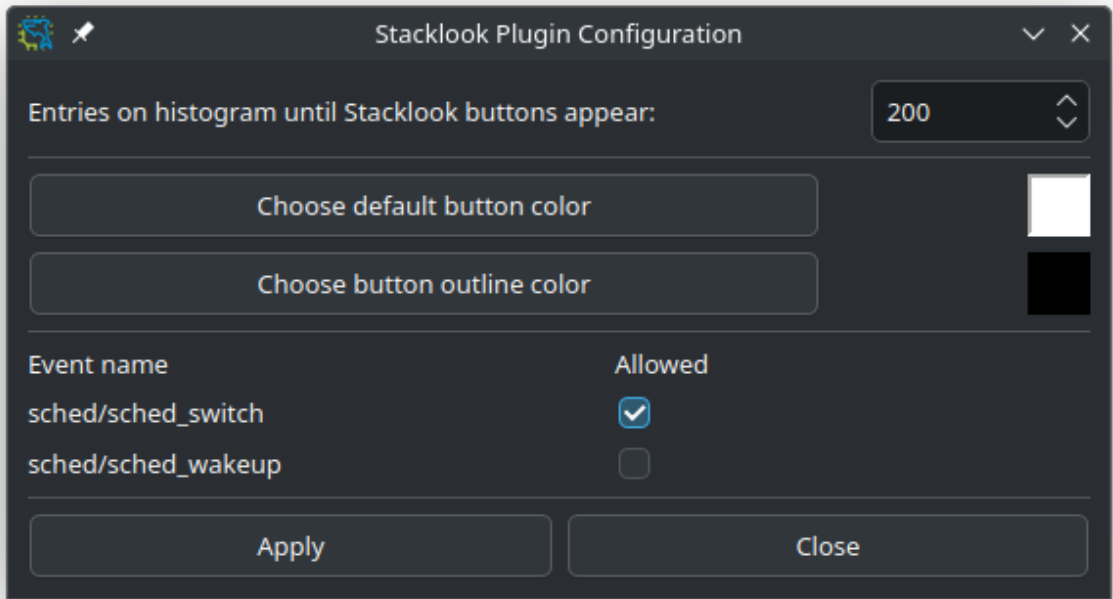
Each setting has an explanation to go along with it, but usual control elements work as expected: check boxes are used for toggling, spin boxes adjust numerical values, buttons with color next to them serve to open a color dialog window and choose a different color. The "Apply" button will save the changes made - if not pressed, any changes made won't take effect. Only saved changes show up in the control elements (e.g. a feature that is turned off will have its check box unchecked). The "Close" button and top-right corner "X" button will close the dialog without additional actions.

Settings are session-specific, so to remember them, use the session export/import feature of KernelShark.

With everything enabled:



With unmodified KernelShark specified and nap rectangles disabled:



Use in a plot

To use the plugin, have a trace file ready which includes *sched/sched_switch* and *sched/sched_wakeup* event entries.

After loading the plugin, zoom until less than the configured number of visible entries are visible in the histogram plot. If there are any of the above entries, a button will show up, colored like the task the event is for.

Since the color theme can be changed in KernelShark, the buttons respect this possibility, albeit they need the user to zoom out and back in for this to take effect.

Excuse the lower resolution 😊



If toggled on via configuration and build, nap rectangles will also show up after the number of entries condition is satisfied. Each of them has outlines the color of the task the *sched/sched_switch* entry is from & their inner color is decided according to the below list:

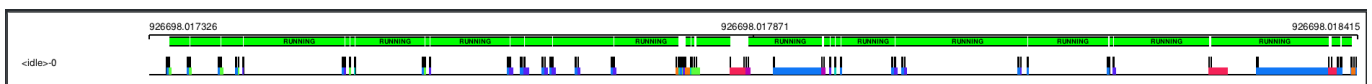
- **S** - sleeping: **blue**
- **D** - uninterruptible (disk) sleep: **red**
- **R** - running: **green**
- **I** - idle: **yellow**
- **T** - stopped: **cyan**
- **t** - tracing stop: **brown**
- **X** - dead: **magenta**
- **Z** - zombie: **purple**
- **P** - parked: **orange**

How a nap rectangle looks like:



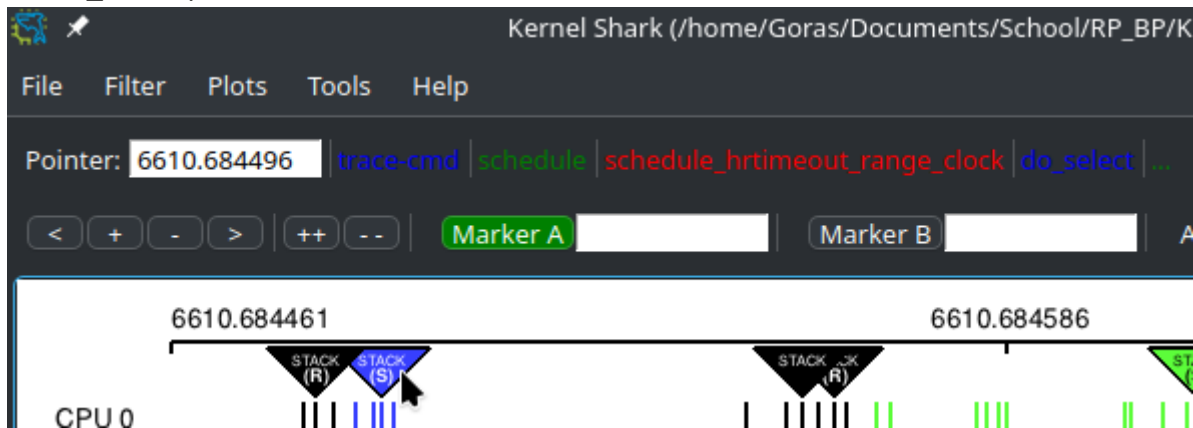
If *ftrace/kernel_stack* events were not collected, no Stacklook buttons will be shown, but nap rectangles will be displayed even if no stack information has been collected, as they don't at all rely on it.

As visible here:

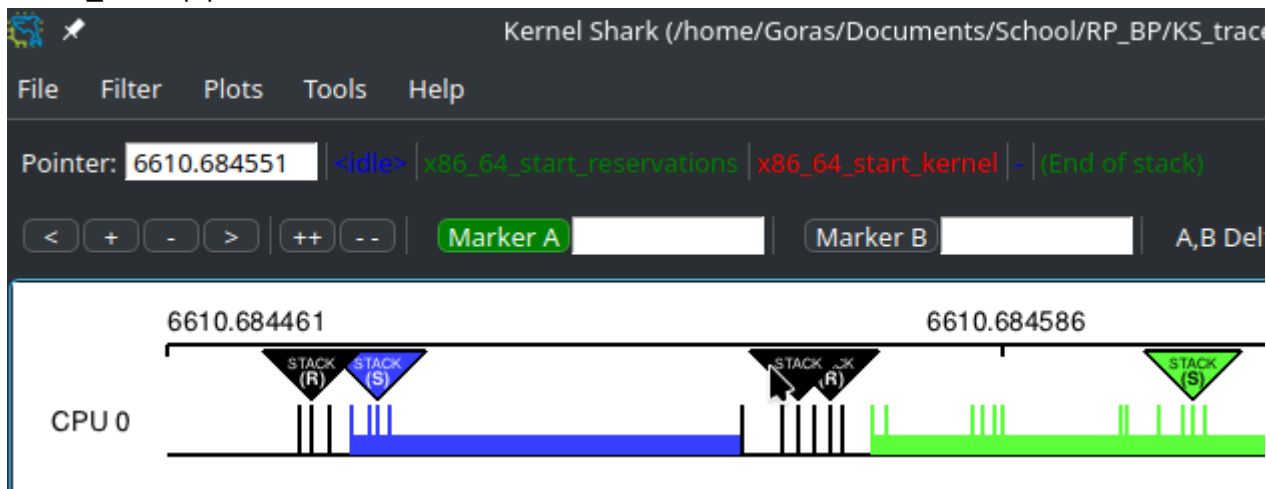


Else, upon mouse hover/move over a button, the name of the task and three items from the top of the kernel stack will be shown. These three items can be configured to start from the very first stack item or they can be offset in the configuration by a certain amount. If there are no entries to display (or we are going over the stack), dashes will be shown instead and the last label in the preview will notify us that we are seeing the end of the stack.

Sched_switch preview, offset is set to 3:



Sched_wakeup preview, offset is set to reach end of kernel stack:

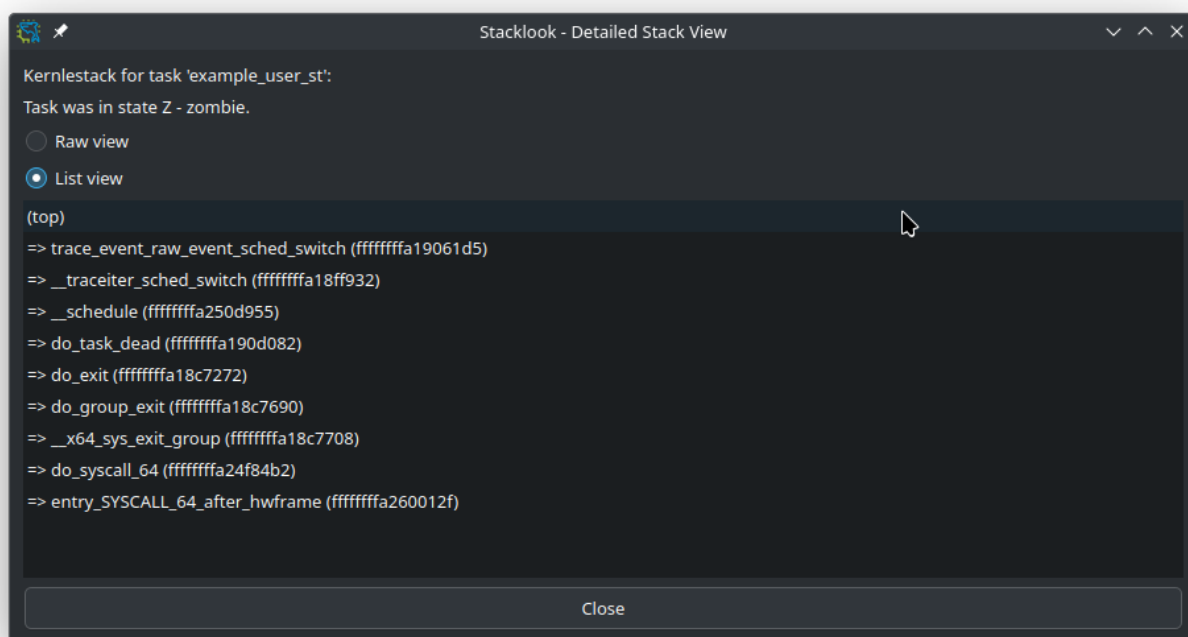


Moving away from the button's boundaries will show nothing, i.e. the labels will be cleared.

Double-clicking on a button will bring up a new window, which will show the full kernel stack, either as raw text or in the form of a list of strings. Both will have "(top)" written above the stack, signifying where the top of the stack is. It is possible to toggle between the raw text view and the list view using two radio buttons above the view area. On the top of the window is a small message with the name of the task from which we had the kernel stack traced. Right below it is a message specific to an event - sched_switches get information about their prev_state, sched_wakeups just reiterate that the task has woken up.

The window can be closed with the "Close" button at the very bottom of the window or with the "X" button of the window's header. Last option is to close the main KernelShark window, which will close all of Stacklook's opened windows.

There can be more than one window opened for a single event and there can be more windows with different events open (performance shouldn't be demanding, but it isn't recommended to open hundreds upon hundreds).



Using Stacklook as a library

It's not really a library

Stacklook was built as a standalone piece of software, not meant as a library (it is a shared library solely because this is what KernelShark expects).

Still, since there are header files, it is good to mention what is in them just in case it is needed.

What's in the header files?

I recommend looking into the **technical Doxygen documentation**. Below are only short summaries of each member of header files.

In **stacklook.h**:

- Plugin context struct `plugin_stacklook_ctx`, which serves as a sort of plugin-wide global variable collection. It stores event IDs of collected events, container of collected events and a pointer to a collection of plugin's opened windows.

A possible extension could be a collection of plugin's buttons so that one can manipulate them from anywhere.

- Global functions:
 - `get_font_ptr` - getting a pointer to KernelShark's font
 - `draw_plot_buttons` - draw handler for drawing buttons on the plot
 - `plugin_set_gui_ptr` - gets a pointer to KernelShark's main window during plugin load
 - `clean_opened_views` - cleans the opened window collection
 - `init_views` - initializes the opened window collection
 - Last four are defined in C++, first one in C. This is due to the necessities of using all of them in C code, but them needing to work with C++ structures.

In `SButton.hpp`:

- Class `SITriangleButton`, used for drawing the buttons. It is also responsible for the necessity of modified KernelShark build, due to it having a mouse hover action defined. It is also used to draw the buttons in a way which makes their parts not bleed into each other, thanks to it being a composite of other plot objects.

This class also houses all logic of plugin's buttons' inner workings. As a child of the general plot object it can be used in graphs. In fact, it is, without this plugin's overall context, just a button that can do actions like a plot object and stores which entry it is gaining data from.

In `SIDetailedView.hpp`:

- Class `SIDetailedView`, class that represents a window with more detail about a kernel stack entry of an event. Class itself holds a pointer to the main KernelShark window and a pointer to all opened windows. Otherwise it is mostly composed of other Qt objects.

In `SConfig.hpp`:

- Usings:

- `allowed_t` - boolean for if an event is allowed to have a button displayed above its entries
- `event_name_t` - string for event names
- `events_meta_t` - map of names to plugin meta information, which depends on whether `_UNMODIFIED_KSHARK` is set or not
- Usings in effect only if `_UNMODIFIED_KSHARK` is NOT set:
 - `depth_t` - integer specifying offset into the kernel stack, used for the preview
 - `event_meta_t` - pair of `allowed_t` and `depth_t`, creating the meta type
- Class `SlConfig`, plugin configuration singleton object. Houses configurable values, but only lets them be read. Changes to it can be made only via `SlConfigWindow` in the GUI.
- Class `SlConfigWindow`, Qt widget, child of the main window. Provides GUI elements to manipulate the `SlConfig` object.

In `SIPrevState.hpp`:

- Static constant map `LETTER_TO_NAME` which uses `prev_state` single-letter abbreviated names as keys for the full name.
- Global functions:
 - `get_switch_prev_state` - returns a single character-long C++ string with the abbreviated name of the `prev_state` of the `sched_switch` event entry in the argument
 - `get_longer_prev_state` - returns a C++ string with the full name of the `prev_state` of the `sched_switch` event entry in the argument

In `SINapRectangle.hpp`:

- Class `SlNapRectangle`, a `KernelShark` plot object used for drawing nap rectangles. Nap rectangles cannot be interacted with and should only be drawn as an interval plot between a `sched_switch` and `sched_wakeup` events.